Screen Space Fluid Rendering with Curvature Flow

Maarten Terpstra* and Bram Musters*

*Department of Computing Science
University of Groningen
Nijenborgh 9
Groningen 9747 AG
{m.l.terpstra, b.t.musters}@student.rug.nl

Abstract

Fluids can be simulated using Smoothed Particle Hydrodynamics(SPH). However, visualizing these fluids in a realistic way proves to be hard. The paper by van der Laan et al. [2009] proposes a method that visualizes this SPH simulation as a nature-like fluid using surface smoothing by screen-space curvature flow. The advantage of this method is that it can be rendered in real-time for a high amount of particles without a need for a finite grid. This paper analyzes the proposed method and describes how to implement it using C++, OpenGL and the nVidia Cg toolkit.

Keywords: Smoothed Particle Hydrodynamics, fluid, visualization, computer graphics

1 Introduction

For some applications it may be useful to simulate fluids using computers in order to improve usage of the fluids in certain situations. By using visualizations of fluids, bottlenecks or missed opportunities in systems can be identified. This kind of visualizations are mainly used in engineering.

Other types of visualizations, which are mainly used in academia, includes fields of research as oceanography, volcanology and astrophysics. In these kind of visualizations, entire systems are modeled to identify breaking points or how different parameters may affect an entire system.

There are several methods available for simulating fluids, such as the Eulerian method or the Smoothed Particle Hydrodynamics (SPH). The Eulerian method is fairly straightforward, as it describes fluid motion as an integration of the surface over time. This simple technique comes with a downside in the fact that it requires a (finite) grid on which the fluid moves, which is computationally expensive and limits the fluid in that it cannot move everywhere. SPH proves to be a viable alternative. With this technique, a body of fluid is reduced to a number of particles, where each particle has a position, velocity, acceleration and density. These properties can be used to extract a smoothed surface which can then be visualized in an aesthetically pleasing way.

Simulating fluids without a visualization may not achieve the desired insight in the system. The insight gained from a fluid rendering can be increased by having a natural visualization of the system.

2 Problem Definition

The goal of this project is to provide a realistic and real-time visualization of fluids modelled by SPH. A regular SPH simulation can be observed in figure 1 and it is hard to argue that it appears to be a fluid. As this definition of the problem is fairly abstract, we will try to specify it further for the intended goal.

The visualization part of the definition is the

Figure 1: An SPH simulation

most obvious term for further specification. In this context it means that the simulation must end up appearing the same as the fluid in nature. This means that a fluid has a color, reflects incoming light and refracts light. The combination of reflection and refraction result in a blending of the fluid with the background. Moreover, the fluid surface must appear to be smooth. In nature, this effect is also visible due to surface tension. It is not a trivial task to quantify the level of realism of a fluid visualization but correct lighting, color, and surface appearance are essential for a realistic fluid visualization. There are other properties to make a fluid appear realistic such as foam, small surface distortions, and "thickness" but do not add as much to the immersion as the aforementioned properties.

The paper by van der Laan et al. van der Laan et al. [2009] achieves a real-time visualization by a point-splatting approach with a configurable speed-versus-quality trade-off. They state that this approach is preferable over methods as implicit surface polygonization and other mesh-like constructs due inherent speed limitation.

van der Laan et al. achieve smooth fluid surfaces by applying a sophisticated smoothing method to the point splats which is discussed in detail in section 3.2.

2.1 Related work

Various methods are developed that try to achieve a natural rendering of fluid.

A method to render fluids has been proposed by Williams [2008] that uses Marching Tiles. The method creates smooth surfaces but the drawback is that these surfaces can not be rendered in real-time.

An improvement Rosenberg and Birdwell [2008] proposed is to extract the isosurface to reduce the amount of particles that need to be rendered. The Marching Cubes algorithm is used as a base for this extraction technique. The advantage is that rendering becomes relatively fast, however the mesh is not smooth and since it is rendered directly, post-processing the result can not be done.

The problem of both Williams [2008] his method and Rosenberg and Birdwell [2008] their method is that they require a fixed grid, which is not desirable for our implementation.

Müller et al. [2007] proposed a method where the boundary of a mesh is generated, without the need of a fixed grid. The surface is smoothed using a binomial filter and creates an intermediate mesh. The drawback is that this method is computationally very expensive, thus this method can not render water in real-time.

Zhang et al. [2008] developed a method that makes use of point-based rendering, therefore a grid is unnecessary. However, a drawback of this method is that it results in unreasonably thick surfaces.

3 Solution

The paper introduced by van der Laan et al. [2009] provides a fitting solution to our problem. The paper describes a method for visualizing fluids simulated using SPH in a natural way. According to the paper, they want to:

1. Achieve real-time performance, with a configurable speed versus quality trade-off.

2. Do all the processing, rendering and shading steps directly on the graphics hardware.

3. Smooth the surface to prevent the fluid from looking blobby or jelly-like.

2

4. Have a method that is not based on polygonization, thus not suffering from the associated tessellation artifacts.

Items 1 and 3 have a direct connection with our goals and items 2 and 4 are ways to enable these goals.

The method described in the paper consists of multiple passes. It assumes that there is a functional SPH simulation where each fluid particle has a density. Then each frame the following steps are performed:

1. Splat points as spheres and determine depth values per fragment

2. Smooth the spheres based on curvature flow

3. Attenuate colors based on thickness

4. Add noise texture and advect throughout the simulation

5. Add foam

6. Render using Fresnel and Phong equations

3.1   Depth determination

It is desirable to determine depth of particles in the simulation in order to know which particles are part of the surface. In order to obtain these depth values, the particles from the SPH simulation are splatted to discs and their fragments are subjected to a hardware depth test. This result is subsequently written to a depth buffer. Note that merely the depth value is splatted and the color and normal attributes are left unchanged as they will be changed in later steps.

3.2   Surface smoothing

Smoothing is a critical component in the paper. It is responsible for creating a smooth surface from a set of points so to avoid a blobby and jelly-like surface. The paper argues that this is a better approach than using Gaussian blurring because it is computationally less expensive and ought to produce better results, because there will be no silhouette blurring. This is achieved by translating fragments along the z-axis according to the curvature flow in that region.

Curvature flow is defined as the divergence of the unit normal by $2H = \nabla \cdot \hat{\mathbf{n}}$. Subsequently, depth values can be displaced every timestep based on the curvature in that point, as $\frac{\partial z}{\partial t} = H$.

To obtain normals, the point of which the normal needs to be determined is mapped back to a point in view-space by inverting the projection transformation. This point is called

$$\mathbf{P}(x,y) = \begin{pmatrix} \frac{\frac{2x}{V_x}-1.0}{F_x} \\ \frac{\frac{2y}{V_y}-1.0}{F_y} \\ 1 \end{pmatrix} z(x,y) = \begin{pmatrix} W_x \\ W_y \\ 1 \end{pmatrix} z(x,y) \quad (1)$$

where V_x and V_y are the viewport dimensions and F_x and F_y is the focal length in x and y component.

The normal is then obtained by

$$\mathbf{n}(x,y) = \frac{\partial \mathbf{P}}{\partial x} \times \frac{\partial \mathbf{P}}{\partial y} \quad (2)$$

Rigorous algebra shows that

$$\mathbf{n}(x,y) = \begin{pmatrix} -C_y \frac{\partial z}{\partial x} \\ -C_x \frac{\partial z}{\partial y} \\ C_x C_y z \end{pmatrix} z \quad (3)$$

Here $C_x = \frac{2}{V_x F_x}$ and $C_y = \frac{2}{V_y F_y}$

Subsequently, the unit normal is obtained by $\hat{\mathbf{n}}(x,y) = \frac{\mathbf{n}(x,y)}{|\mathbf{n}(x,y)|} = \frac{\mathbf{n}(x,y)}{\sqrt{D}}$, where

$$D = C_y^2 \left(\frac{\partial z}{\partial x}\right)^2 + C_x^2 \left(\frac{\partial z}{\partial y}\right)^2 + C_x^2 C_y^2 z^2 \quad (4)$$

Now that we have a unit normal we can define mean curvature as

$$2H = \frac{\partial \hat{\mathbf{n}}_x}{\partial x} + \frac{\partial \hat{\mathbf{n}}_y}{\partial y} = \frac{C_y E_x + C_x E_y}{D\sqrt{D}} \quad (5)$$

in which

$$E_x = \frac{1}{2} \frac{\partial z}{\partial x} \frac{\partial D}{\partial x} - \frac{\partial^2 z}{\partial x^2} D \quad (6)$$

$$E_y = \frac{1}{2} \frac{\partial z}{\partial y} \frac{\partial D}{\partial y} - \frac{\partial^2 z}{\partial y^2} D \quad (7)$$

These values can be calculated fairly easy and can thus be used to modify the z-values in the visualization quickly using Euler integration in time.

3.3 Thickness

The thickness of a fluid is important because an object behind the fluid should change color according to the amount of fluid in front of it. Also, the thickness is used to determine the transparency and color of the fluid. This means that the distance from the camera towards the first opaque object has to be determined for each fragment.

Each particle is rendered as a sphere and has a fixed size. Additive blending is used to determine the total thickness of a surface at each fragment.

3.4 Noise

In order to remove the appearance of an artificially smooth surface, noise has to be added. Since the simulation has moving particles in it, the noise has to advect with those particles for a realistic result. Also the noise has to have a smaller and larger frequency than the fluid.

van der Laan et al. [2009] suggests to use Perlin [1985] noise. One octave of noise is added to each projected particle, rather than to the entire simulation to make sure that noise advects with the particles. The amount of noise that is added to each particle depends on the depth of the particle with respect to the surface. In other words, noise contributes less as particles submerge.

Equation 8 shows this noise calculation in a formula.

$$I(x,y) = noise(x,y) * e^{-x^2 - y^2 - (\mathbf{p}_z(x,y) - d(x,y))^2}$$
$$(8)$$

, where $\mathbf{P}$ is the position of the pixel in view-space, d the depth, and x and y vary between -1 and 1.

The noise of every particle is then summed together to get the correct amount of noise per fragment. Fluids that have a high change in velocity need to be marked, so the noise for that particles is influenced more. Additive blending is used so the Perlin noise is correctly blended into the fluid.

Also, foam can be added to the fluid by adding grey according to the magnitude of the noise.

3.5 Rendering

Finally, all intermediate results are merged together in a final texture that is rendered as fullscreen quad. Since light changes direction when it moves between media, and thus the reflection and refraction vectors changes too, a calculation for this direction change is needed. A visualization of this behavior can be seen in figure 2.


Figure 2: Light direction change between surfaces (From: http://en.wikipedia.org/wiki/File:Fresnel1.svg)

The Fresnel equation (by Augustin-Jean Fresnel) is used to calculate the correct reflection and refraction vectors,

$$R_s = \left| \frac{n_1 \cos\theta_i - n_2 \cos\theta_t}{n_1 \cos\theta_i + n_2 \cos\theta_t} \right|^2 \qquad (9)$$
$$T_s = 1 - R_s. \qquad (10)$$

Where R_s is the reflection vector, and T_s is the refraction vector. The meaning of the other variables can be seen in figure 2. Once these vectors are known, the Phong [1975] illumination model can be used for the rendering.

This results in the following equation to determine the output color van der Laan et al. [2009]:

$$C_{out} = a(1 - F(\mathbf{n}\cdot\mathbf{v})) + bF(\mathbf{n}\cdot\mathbf{v}) + k_s(\mathbf{n}\cdot\mathbf{h})^\alpha. \quad (11)$$

F is the Fresnel equation, a is the refracted fluid color, b is the reflected scene color, k_s and α are both constant and represent the specular highlight. $\mathbf{n}$ is the normal along the surface, $\mathbf{h}$ is the half-angle between the camera and the light, and $\mathbf{v}$ is the direction pointing towards the camera.

4 Implementation

The solution that has been proposed in the previous section is implemented using the C++ programming language in combination with the OpenGL API. To program on the GPU, Cg is used as shading language.

The actual SPH fluid simulation is given and can be seen in figure 1, by implementing the multiple passes described in the previous section, a nature like fluid is expected. By creating Frame Buffer Objects (FBO), the results can be written to an off-screen rendering target. This is useful since we are using multiple passes. Also, multiple buffers can be attached to the FBO.

4.1 Depth determination

The depth at each fragment of each particle closest to the camera is determined using a combination of a vertex- and fragment shader. Each particle has a position in world space that is passed to the vertex shader. These particles are rendered as spheres to determine the correct depth values.

The vertex shader computes and passes the following properties of each particle to the fragment shader:

- Position of center in eye space.
- Position of center in screen space.
- Splat size.
- Splat radius.

The fragment shader uses this input to determine the depth at every fragment. To determine this depth, the splat is rendered as a sphere by discarding fragments that fall outside the sphere. From this, the normal from the center of the sphere towards its surface is determined by taking the difference of the current fragment position and the current particle center. Using this normal, the point is transformed to clip space, the z value of this position is the depth value. The depth values are then written to the depth buffer of the current FBO.

The depth component of each fragment can be visualized as grey value by setting the R, G and B components to the depth value. The results of this visualization can be seen in figure 3. It can be observed that the particles become darker when they appear closer to the camera.


Figure 3: Visualisation of the depth components

4.2 Surface smoothing

Smoothing is done using a fragment shader. The original depth values are modified using curvature flow to smooth the surface. Since one cannot modify the depth texture that is used for reading, there are two depth textures which are alternated during computation.

The smoothing is determined in multiple iterations and depth values are modified in several timesteps. In pseudocode, the algorithm is as follows.

Algorithm 2 Surface smoothing pseudocode
---
for $i := 1$ to smoothSteps do
    smoothDepths();
    flipTextures();

4.2.1 Depth smoothing

Using the depth values, finite differencing can be applied to determine derivatives of these values to obtain the curvature in every point, as equation 5 describes. This curvature is used to modify the depth values which are subsequently written to another depth texture. We had trouble implement-

ing this step in Cg, but we found a reference [1] that implemented the smoothing in GLSL, which we adapted to our needs.

4.3 Rendering

For the final rendering phase we implemented a fragment shader that produces a color for each fragment, which only has a depth component. We have a default color - RGB(0.02, 0.1, 0.3), with $\alpha = 0.5$ - for water that is modified according to its surroundings based on shading. The end-results are projected onto a full-screen quad, which is the best option due to our framebuffer-objects.

The paper suggests to implement a combination of Phong shading and the Fresnel equation for displaying the surface. Normally, computing the Fresnel equation would be a costly and difficult operation. To circumvent this, we have applied Schlick's approximation to determine the reflective Fresnel component Schlick [1994].

4.3.1 Normal calculation

To determine the normals of the surface, the depth buffer is passed as input to the fragment shader, along with the uniform variables C_x and C_y. Since the depth buffer is in a texture, we can obtain the depth values of neighbors. Using these neighbors, the finite differences between depth values can be calculated in the x and y direction. Since we now have $C_x, C_y, \frac{\partial z}{\partial x}, \frac{\partial z}{\partial y}$ and z, we can approximate the normals according to equation 3. A visualization of the normals, done by assigning the normal to the color output, can be seen in figure 4 and 5. These normals are essential for a correct shading of the surface, but are not explicitly needed for the smoothing of the surface.

5 Conclusion

When the initial requirements for the realistic water rendering are considered it seems that the curvature flow way is a good fit.

- Given enough iterations of the fluid surface, a surface will become smooth and will have no sharp discontinuities between particles.


Figure 4: Visualization of the normals of a unsmoothed surface

Figure 5: Visualization of the normals of a smoothed surface

- It is possible to generate an imperfect surface by adding a noise texture to the surface and attaching it to particles. This will distort the

---

[1] https://github.com/halcy/simpleflow/blob/master/WaterSim2/curvatureflow.frag

surface a little bit and will create an imperfect, and more natural surface

- It takes into account attenuation of colors based on the depth of the fluids and submerged objects. Moreover, foam can be added for an added realistic effect.

Besides that it can produce a natural rendering of a fluid surface, it is also fast. It can render a smooth surface from thousands of particles and still have an acceptable framerate. This is also an important quality for a simulation.

5.1 Results

We have implemented the methods as described in the Implementation section. The results are visible in figure 6 and 7.


Figure 7: The same simulation visualized with 24 smooth steps

Figure 6: The SPH simulation with splatted point and Phong and Fresnel rendering without surface smoothing

cused more on the solution as proposed by the paper and the introduction.

All in all, we feel that we both contributed equally to the entire final project.

References

Müller, M., Schirm, S., and Duthaler, S. (2007). Screen space meshes. In *Proceedings of the 2007 ACM SIGGRAPH/Eurographics symposium on Computer animation*, 9–15. Eurographics Association.

Perlin, K. (1985). An image synthesizer. *ACM Siggraph Computer Graphics*, **19**(3), 287–296.

Phong, B. T. (1975). Illumination for computer generated pictures. *Communications of the ACM*, **18**(6), 311–317.

Rosenberg, I. D. and Birdwell, K. (2008). Real-time particle isosurface extraction. In *Proceedings of the 2008 symposium on Interactive 3D graphics and games*, 35–43. ACM.

Schlick, C. (1994). An inexpensive brdf model for physically-based rendering. In *Computer graphics forum*, vol. 13, 233–246. Wiley Online Library.

6 Workload division

All the programming was done in pairs so we'd value that the programming was 50-50 effort.

For the report, Bram focused more on problem definition, and implementation while Maarten fo-

van der Laan, W. J., Green, S., and Sainz, M. (2009). Screen space fluid rendering with curvature flow. In *Proceedings of the 2009 symposium on Interactive 3D graphics and games*, 91–98. ACM.

Williams, B. W. (2008). *Fluid surface reconstruction from particles*. Ph.D. thesis, Citeseer.

Zhang, Y., Solenthaler, B., and Pajarola, R. (2008). Adaptive sampling and rendering of fluids on the gpu. In *Proceedings of the Fifth Eurographics/IEEE VGTC conference on Point-Based Graphics*, 137–146. Eurographics Association.